



Faculty of Mathematics and Information Science  
Warsaw University of Technology



# Multigame Playing - A Challenge for Computational Intelligence

Jacek Mańdziuk



# Grand challenges - MULTIGAME PLAYING



- Deep Blue II is unable to play simple tic-tac-toe
- **Meta-level approach** (autonomous game analysis, learning and problem solving)
- Focus on universal, generic AI/CI algorithms rather than solutions and methods optimized for single (particular) game:
  - Abstract reasoning
  - Knowledge representation
  - Game-independent learning
  - Planning
  - Knowledge transfer
  - Life-long learning
  - ...
- **Return to the idea of Strong AI (AGI)**

# Multigame playing - SAL



- Search And Learning (M. Gherrity, 1993)
- Deterministic, two-player, perfect-information board games
- Game independent kernel with the general knowledge about this game genre – pre-defined, unchanged
- Game related knowledge (move-generator, constants defining the board, the pieces, etc.) provided by the user in a separate game-specific module (written in C)
- Two evaluation functions - 1hl MLP (→ non-zero-sum games)
- Each represented as a NN with TD-learning
- Positional features (pieces, quantities), Dynamic features (moves, captures) and rule-based features (threatened pieces, controlled squares, etc)
- Shallow search (2-ply only)
- EXTREMELY SLOW LEARNING
- EXPLOSIVE NUMBER OF FEATURES

# Multigame playing - HOYLE



- **Hoyle** (Susan Epstein, 1991-2004)
- Two-player, perfect-information, deterministic, finite board games
- Only the rules of the game are provided
- Three tiers of advisors commenting on particular aspects of a game position
- Tier 1: binary assessment of particular moves from the point of view of a given Advisor, e.g. Victory Advisor – winning moves, Wiser – winning in 1 ply, Sadder Advisor – losing, Don't Lose – not losing in 1 ply, etc.)
- Tier 2: advocating certain plans of play (achieving particular goals)
- Tier 3: *non-binary* assessment of moves (in Advisor's expertise area) – Fork, Freedom (Mobility), Material, etc.
- Tiers 1 and 2 – sequential decision. If not conclusive then weighted decision made in Tier 3
- **REASONABLE LEVEL OF PLAY** in 18 two-player games (incl. t-t-t).
- **WEAK CHESS playing (novice level)**

# Multigame playing - METAGAMER



- METAGAMER (Barney Pell, 1992-1996)
- Any „symmetric chess-like” (SCL) game, i.e. chess, checkers, shogi
- Universal evaluation function – linear combination of simple pre-defined features (partial goals)
- Similarly to Hoyle, some number of Advisors was employed, which voted on potential usefulness of particular features (in the form of numerical estimation of the worth of particular features)
- Among 23 Advisors there were 4 related to *mobility*, 4 to *threats and capturing*, 4 to *goal functions* and 11 to *material values*
- Shallow search – 1 ply with 1 ply non-quietness search
- Novice level in chess (GNU Chess level 1 with a Knight handicap) and intermediate level in checkers
- No positive results in ad-hoc defined SCL games



# PART 2

## General Game Playing UCT search

# General Game Playing



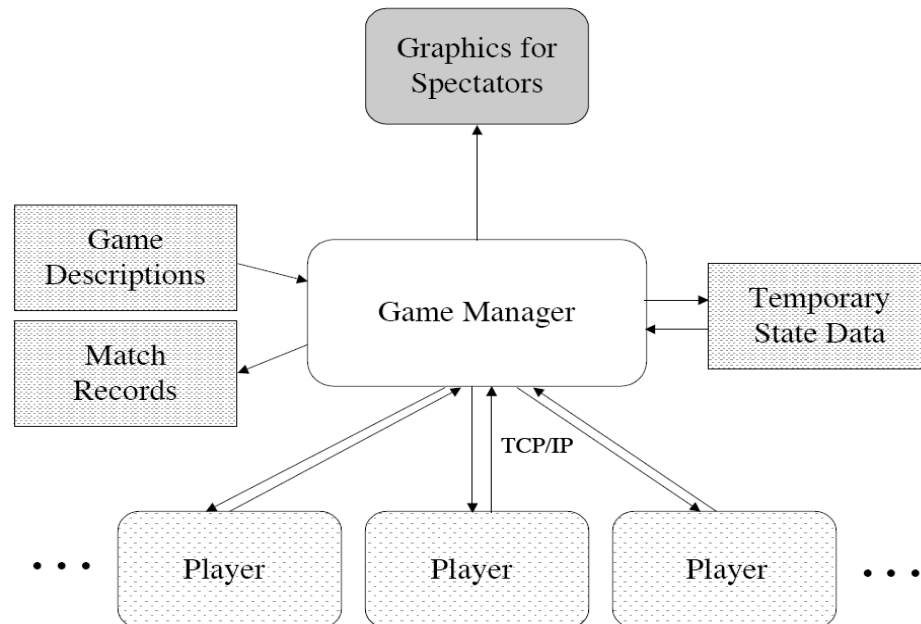
- Building programs (agents) able to play any game (within a certain class of games) without human intervention.
- Annual GGP tournament (since 2005)
- Stanford online course (2014)
- Game controlled by a GameMaster
- Communication via HTTP

# General Game Playing



- GameMaster
  - Provides game description to each player
  - Verifies legality of proposed actions
  - Sends out info about (other) players' actions
  - Monitors time limits

- Time limits
  - StartClock
  - PlayClock



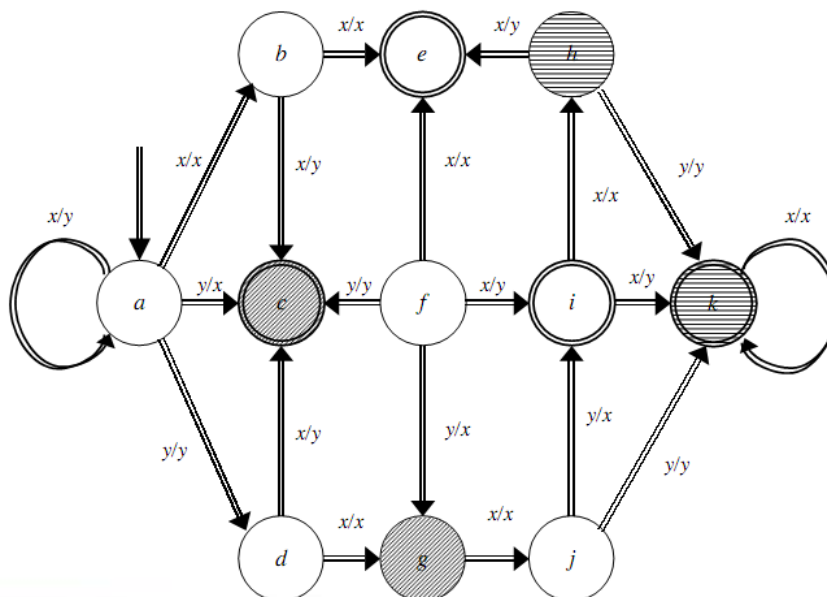


# Genre of games



Jacek Mańdziuk, Metody Sztucznej Inteligencji 2

- Multiplayer
- Finite
- Synchronous
- Deterministic



M. Genesereth, N. Love, and B. Pell. General Game Playing: Overview of the AAAI Competition. AI Magazine, 26(2):62-72, 2005.

# Game Definition Language (GDL)



- Allows more concise description of game rules
- Uses a version of Datalog – a subset of Prolog
- Specific notation of variables, relations, logic operators
- A set of keywords

# Main GDL keywords/relations



(role xplayer)  
(role oplayer)

(<= (column ?n ?x)  
(true (cell 1 ?n ?x))  
(true (cell 2 ?n ?x))  
(true (cell 3 ?n ?x)))

(<= (goal xplayer 100)  
(line x))

(init (cell 1 1 b))  
(init (cell 1 2 b))  
...  
(init (cell 3 3 b))  
(init (control xplayer))

(<= (diagonal ?x)  
(true (cell 1 1 ?x))  
(true (cell 2 2 ?x))  
(true (cell 3 3 ?x)))

(<= (goal xplayer 50)  
(not (line x))  
(not (line o))  
(not open))

;; Cell  
(<= (next (cell ?m ?n x))  
(does xplayer (mark ?m ?n))  
(true (cell ?m ?n b)))

(<= (diagonal ?x)  
(true (cell 1 3 ?x))  
(true (cell 2 2 ?x))  
(true (cell 3 1 ?x)))

(<= (goal xplayer o)  
(line o))  
  
(<= (goal oplayer 100)  
(line o))

(<= (next (cell ?m ?n o))  
(does oplayer (mark ?m ?n))  
(true (cell ?m ?n b)))

(<= (goal oplayer 50)  
(not (line x))  
(not (line o))  
(not open))

...

(<= (line ?x) (row ?m ?x))  
(<= (line ?x) (column ?m ?x))  
(<= (line ?x) (diagonal ?x))

(<= (goal oplayer o)  
(line x))

(<= (next (control xplayer))  
(true (control oplayer)))

(<= open  
(true (cell ?m ?n b)))

(<= (next (control oplayer))  
(true (control xplayer)))

(<= terminal  
(line x))

(<= (row ?m ?x)  
(true (cell ?m 1 ?x))  
(true (cell ?m 2 ?x))  
(true (cell ?m 3 ?x)))

(<= (legal ?w (mark ?x ?y))  
(true (cell ?x ?y b))  
(true (control ?w)))

(<= terminal  
(line o))

(<= (legal xplayer noop)  
(true (control oplayer)))

(<= terminal  
(not open))

(<= (legal oplayer noop)  
(true (control xplayer)))

**role(r)**  
**init(p)**  
**true(p)**  
**legal (r,a)**  
**does(r,a)**  
**next(p)**  
**terminal**  
**goal(r,value)**

# GGP simulation scheme



- Calculate legal actions [legal]
- For each role  $i, i = 1, \dots, N$  [does]
  - move[i] (select an action)
  - perform(move[i]) (send out all moves)
- Calculate the next game step [next]
- Check if terminal [terminal]
  - YES  $\rightarrow$  goal(role,value) and STOP [goal]
  - No  $\rightarrow$  goto [legal]

# Games Played

GGP Player



## Chess



IBM Deep Blue II

Chess  
Checkers  
Chinese Checkers (3,4,6)  
Connect Four  
Connect Five  
Connect Four Suicide  
Othello  
Quarto  
Pentago  
Pilgrimige  
Blocker  
Tic-Tac-Toe  
9-Board Tic-Tac-Toe  
Numeric Tic-Tac-Toe  
Tic-Tac-Chess (3,4)  
Counterstrike (Simplified)  
Eight Puzzle (solving)



Sudoku Puzzle (solving)  
Breakthrough  
Knight-through  
Free For All (2,3,4)  
Rock-Paper-Scissors  
Lights Out  
Cephalopod  
Chess and Othello  
(simultaneously)  
Farmers (3)  
Zhadu  
Nim  
Cylinder Checkers  
Nine Men's Morris  
Finding a Knight's Tour  
Flipping Pancakes  
...

# GGP approaches - summary



- 2005 – 2006
  - heuristic-based syntactical game analysis
  - construction of the evaluation function
- 2007-2016 → **MCTS** (cont. GVGP)
  - simulation-based methods
  - Cadia Player (2007, 2008, 2012) – Reykjavik University, Iceland
  - Ary (2009, 2010) – Paris 8, France
  - TurboTurtle (2011, 2013) – USA
  - Sancho (2014) – USA & UK
  - Galvanise (2015)
  - WoodStock (2016) - France

# MCTS - MOTIVATION



- MCTS combines the focused and precisely-defined tree search with the power of massive random sampling of the problem space
- MCTS can be applied to any problem which relies on sequential decisions implemented as a tree-search process (whose solutions can be represented as paths in a tree)
- MCTS is a well suited method for problems intractable by min-max / alpha-beta search (deep tree, large branching factor, lack of compact and reliable assessment function )
- DAG

# MONTE CARLO METHODS - BASICS



Rooted in statistical physics – numerical approximation of complex integrals.

Used for estimating the quality of actions in certain domains (games, planning) based on **uniform sampling** (now called *flat Monte Carlo*)

**Flat MC** in many domains / problems is ineffective as it does not assume modeling the opponent in any way

Significant improvement in 2006 → UCB1 action selection policy → UCT search tree method



# UCB1 (K-ARMED BANDIT)



- UCB = Upper Confidence Bounds (UCT = UCB applied to Trees)
- k-Armed Bandit Problem or k (one-arm) Bandits Problem
- Distributions of pay-offs  $\{X_{j,t}\}_{t=1,2,\dots} j=1,2,\dots,k$  are fixed, but unknown
- How to maximize the total (long-term) reward?

- Exploration vs. exploitation balance
- First, play each arm once
- Then, use the following rule:



$$A^* = \arg \max_{i=1,\dots,k} \left\{ \bar{X}_i + C \sqrt{\frac{\ln n}{n_i}} \right\}, \quad C = \sqrt{2}$$

# UCT - MULTIPLE SIMULATIONS



Perform multiple simulations according to the following pattern:

Choose an action not yet selected (if exists)

If all had already been tried select action  $a^*$

$$a^* = \operatorname{argmax}_{a \in A(s)} \left\{ Q(s, a) + C \sqrt{\frac{\ln N(s)}{N(s, a)}} \right\}$$

$Q(s, a)$  – the average result for a pair (*state*, *action*)

$N(s)$  – the number of visits in state  $s$

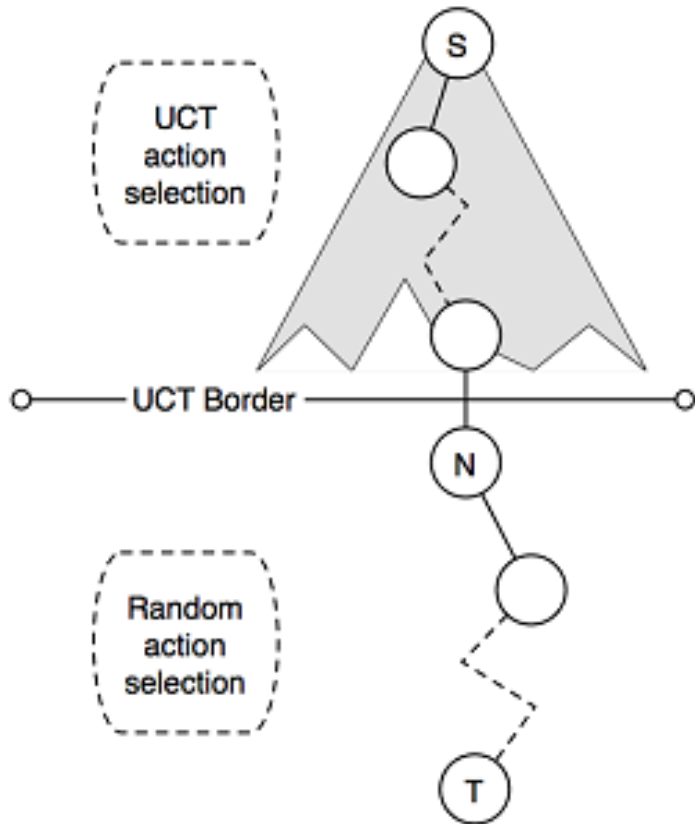
$N(s, a)$  – the number of times action  $a$  was selected in state  $s$

Once the simulation is completed propagate the result back in the tree

**The saliency of  $C$  parameter**

The UCT-based GGP player performs as many such simulations as possible (within allotted time), thus estimating the  $Q(s, a)$

# MCTS WITH UCT



- Due to memory limitations:
  - A game tree is built gradually – one new node is added in each simulation.
  - Once the real (not simulated) action is selected by the system part of the tree above a given node is deleted from memory
- Random simulation to the leaf node.
- Node's assessment is equal to the average payoff of the simulations made so-far.

Bjornsson, Finnsson, IEEE TCIAIG, 1(1), 2009

# MCTS/UCT








selection

expansion

simulation

backpropagation



Tree policy  
(UCB1)

Default policy  
(random sampling)

# MCTS/UCT



- In practice there are some constraints:
  - Computational budget
  - Time constraints
  - Memory limitations
- The system performs as many simulations as possible estimating the  $Q(s,a)$ , in particular  $Q(\text{root}, a)$
- Afterwards the action to be taken (in real NOT simulated mode) is selected

# MCTS - ACTION SELECTION CRITERIA



$\arg \max_a Q(\text{root}, a)$  – *Max child policy*

$\arg \max_a N(\text{root}, a)$  – *Robust child policy*

Continue sampling until both measures point the same node - ***Max Robust child policy***

A combination of *Max* and *Robust* – ***Secure child policy***:

$$\arg \max_a \left[ Q(\text{root}, a) + \frac{k}{\sqrt{N(\text{root}, a)}} \right]$$


# MCTS - HIGHLIGHTS



- Estimates true min-max value
- Anytime
- Aheuristic (Knowledge-free)
- Asymmetric (Best-First-like search)
- Well scaling and easily parallelized




# MCTS - EASY PARALLELIZATION




**Root Parallelization**

**Hybrid Parallelization**



**Leaf Parallelization**



**Tree Parallelization**



# MCTS/UCT







selection

expansion

simulation

backpropagation



Tree policy  
(UCB<sub>1</sub>)

Default policy  
(random sampling)

# TREE POLICY ENHANCEMENTS



- Three types of approaches:
  - Replacing UCB1 by another exploration-exploitation formula → bayesian inference
  - Tuning UCB1 (variations)
  - UCB1 supported with an evaluation function

# UCB1-TUNED



- Distributions of pay-offs  $\{X_{j,t}\}_{t=1,2,\dots}$   $j=1,2,\dots,k$  are fixed, but unknown
- First, play each arm once
- Then, use the following rule:  $A^* = \arg \max_{i=1,\dots,k} \left\{ \bar{X}_i + C \sqrt{\frac{\ln n}{n_i}} \right\}, \quad C = \sqrt{2}$

$$C = \sqrt{\min \left\{ \frac{1}{4}, V_i(n_i) \right\}}$$

where

$$V_i(r) = \left( \frac{1}{r} \sum_{m=1}^r X_{i,m}^2 \right) - \bar{X}_{i,r}^2 + \sqrt{\frac{2 \ln n}{r}}$$

# UCB1 AND EVALUATION FUNCTION



- Suppose we have a „light“ evaluation function:
  - Initial sorting of not-yet-selected actions
  - Initial estimation of actions' strengths based on some number of simulations

# EXPANSION ENHANCEMENTS



- Rarely considered in practice
- One may try to optimize the number of new nodes added
  - Hard to estimate
  - Problem dependent
- One may optimize the number of concurrent simulations
  - The information is not updated (accurate)

# BACKPROPAGATION ENHANCEMENTS



- Weighting simulation results
  - Some simulations are more important than the others
  - Newer simulations preferred over the later ones (esp. in the case of changing environment)
  - Shorter simulations are more accurate than the longer ones
  - Each simulation is assigned a positive integer weight  $w$  and is treated as performed  $w$  times
- Decaying Reward
  - Preference for shorter wins over the longer ones
  - A reward value is decayed by a value of  $\gamma$ ,  $0 < \gamma \leq 1$  between each two nodes on the backpropagation path

# DEFAULT POLICY ENHANCEMENTS



The biggest room for potential improvement

Motivation:

- The true min-max estimation, but ...
- In the case of insufficient number of simulations: there is too much randomness in actions' assessment
- The sequences of actions are often „not realistic“

# HISTORY HEURISTICS - GIBBS DISTRIBUTION



- An action which was successful in the past (regardless of the state) is treated as a potential candidate also in a given state.
- After each simulation statistics for each action are updated (the average performance – expected reward)
- Gibbs distribution:

$$\pi_s(a) = \frac{e^{\frac{Q(a)}{\tau}}}{\sum_{b \in A(s)} e^{\frac{Q(b)}{\tau}}} \quad \tau > 0$$

**Roulette wheel or Epsilon-greedy**



# LAST GOOD REPLY POLICY



- Used specifically in games
- Each move is considered a reply to the opponent's last move and treated as successful if the player won that game
- For each move the last successful reply is stored (frequently overridden)
- In the simulation, if the LGR move is legal it will be played, otherwise the default policy will be used

# DEFAULT POLICY ENHANCEMENTS



- Except for the HH (and its variants) it is hard to show a generally improving method unless some domain knowledge is considered
- The use of the „ad-hoc” constructed (**simple and automatically devised**) evaluation function
  - Replacing the whole playout phase (with some predefined probability)
  - Finishing the playout phase before reaching the leaf node (with some predefined probability in each node)  
→ **MAGICIAN**

# GGP Competition 2012



- Two phases of the tournament
- PHASE 1: Seven games played against selected opponents
- The top 8 advanced to the final round (MP was the 5th, 4 wins and 3 losses)
- PHASE 2 – until two losses
- Finally MP placed 7th
- Connect-4
- Cephalopod Micro
- Free for All 2P
- Pentago
- 9 Board Tic-Tac-Toe
- Connect-4 Suicide
- Checkers
- Farming Quandries
- Pilgrimage

# GGP Competition 2013



- Two phases of the tournament
- PHASE 1: many games including 1,2,4-player ones; 5 advanced to phase 2 (but not MP ☹, and not M ☹)
- MP – a bug in single-player games ☹
- Finally MP placed 11th (out of 16)

# GGP Competition 2014



- Three phases of the tournament
- Massive Stanford online course
- PHASE 1: eliminations (various types of games)
- The top 17 advanced to the second PHASE
- PHASE 2 – top 12 advanced to the final phase
- PHASE<sub>3</sub> – we ended up at the 7-8th place (out of 49 participants)

# Conclusions



- **Multigame playing is besides intuitive playing and creativity one of the hallmarks of human intelligence (in game domain)**
- GGP directly addresses this issue and stimulates research towards accomplishing this goal
- There are some inefficiencies (especially slowness of analysis) which stem from logic description → how to obey/alleviate them?
- GDL (GGP) is potentially extendable to non-deterministic games
- **If we leave out time constraints more established AI learning approaches (neural nets, genetic/memetic methods) can come into play**
- **Applicability of UCT extends beyond game domain (problems with pay-off defined only in terminal states and/or varying in time)**
- Strategy switching mechanism allows for greater flexibility of the method, its self-improvement and adaptability